

The Implementation of the Cilk-5 Multithreaded Language

M. Frigo, C. E. Leiserson, K. H. Randall

PWL: Zurich

Kornilios Kourtis

27 April, 2017

Previously on PWL: Zurich



Today: Cilk-5

- ▶ A parallel programming language
- ▶ how to write parallel programs
- ▶ how to efficiently execute them
- ▶ targets single (shared-memory) machines
- ▶ original paper written in 1998

Why this paper?

- ▶ theory + practice
- ▶ topics: parallel programming, synchronization, languages




BACK TO 1998...

back to 1998...

a linguistically simple “inlet” mechanism for nondeterministic control. Cilk-5 is designed to run efficiently on contemporary symmetric multiprocessors (SMP's), which feature hardware support for shared memory. We have coded many applications in Cilk, including the \star Socrates and Cilkchess

back to 1998...



back to 1998...

"Gettin' Jiggy wit It"



Single by Will Smith

from the album *Big Willie Style*

B-side	"Big Willie Style"
Released	January 27, 1998
Format	CD single, Cassette
Recorded	1997

What followed


- ▶ 2006: Intel releases its first dual-core CPU
- ▶ 2006: Cilk Arts commercializes Cilk
- ▶ 2008: Cilk++ ships
- ▶ 2008: Most influential PLDI paper award for 2008
- ▶ 2009: acquired by Intel
- ▶ 2010: CilkPlus ships with Intel compiler
- ▶ 2014: CilkPlus becomes part of gcc (`-fcilk`)

- ▶ Programming model also picked up by other languages: See Fork/Join in Java <http://www.oracle.com/technetwork/articles/java/fork-join-422606.html>, and OpenMP tasks.

Programming model

- ▶ Task parallelism
- ▶ Extends C
- ▶ two basic keywords: `spawn` and `sync`.


```
/* Cilk example */  
x = spawn A();  
y = spawn B();  
C();
```



Programming model

- ▶ Task parallelism
- ▶ Extends C
- ▶ two basic keywords: `spawn` and `sync`.

```
/* Cilk example */  
x = spawn A();  
y = spawn B();  
C();  
sync;  
/* x,y are available */
```



Example: run-length encoding (RLE)

(nope, not fib)

IN: [a,a,a,a,b,b,b,c,c,c,c,c]

OUT: [(a,4),(b,3),(c,5)]

Sequential RLE

(accumulator)

```
def rle(xs):  
    ret, curr, freq = ([], xs[0], 1)  
    for item in xs[1:]:  
        if item == curr:  
            freq += 1  
        else:  
            ret.append((curr, freq))  
            curr, freq = (item, 1)  
    ret.append((curr, freq))  
    return ret
```

Recursive RLE

```
def rle_rec(xs):  
    if len(xs) <= 1:  
        return [(xs[0], 1)]  
  
    mid = len(xs) // 2  
    rle1 = rle_rec(xs[:mid])  
    rle2 = rle_rec(xs[mid:])  
  
    return rle_merge(rle1, rle2)  
  
def rle_merge(rle1, rle2):  
    if rle1[-1][0] == rle2[0][0]:  
        r1, rle1 = rle1[-1], rle1[:-1]  
        r2, rle2 = rle2[0], rle2[1:]  
        rle1.append((r1[0], r1[1] + r2[1]))  
    return rle1 + rle2
```

Recursive parallel RLE

```
def rle_rec(xs):
    if len(xs) <= 1:
        return [(xs[0], 1)]

    mid = len(xs) // 2
    rle1 = spawn rle_rec(xs[:mid])
    rle2 = spawn rle_rec(xs[mid:])
    sync
    return rle_merge(rle1, rle2)

def rle_merge(rle1, rle2):
    if rle1[-1][0] == rle2[0][0]:
        r1, rle1 = rle1[-1], rle1[:-1]
        r2, rle2 = rle2[0], rle2[1:]
        rle1.append((r1[0], r1[1] + r2[1]))
    return rle1 + rle2
```

Recursive parallel RLE

```
def rle_rec(xs):  
    if len(xs) <= 1:  
        return [(xs[0],
```

NOTE: Removing the spawn and sync keywords results in a valid sequential program. In Cilk, this is called the **C elision**.

```
        mid = len(xs) // 2  
        rle1 = spawn rle_rec(xs[:mid])  
        rle2 = spawn rle_rec(xs[mid:])  
    sync  
    return rle_merge(rle1, rle2)
```

```
def rle_merge(rle1, rle2):  
    if rle1[-1][0] == rle2[0][0]:  
        r1, rle1 = rle1[-1], rle1[:-1]  
        r2, rle2 = rle2[0], rle2[1:]  
        rle1.append((r1[0], r1[1] + r2[1]))  
    return rle1 + rle2
```


product using spawn/sync

```
cilk int product(int *A, int n) {  
    if (n == 1)  
        return A[0];  
    x1 = spawn product(A,n/2);  
    x2 = spawn product(A+n/2,n-n/2);  
    sync;  
    return (x1*x2);  
}
```

also: `inlet` and `abort`

```
cilk int product(int *A, int n) {  
    int p=1;  
    inlet void mult(int x) {  
        p *= x;  
        if (p == 0)  
            abort;  
        return;  
    }  
    if (n == 1)  
        return A[0];  
    mult( spawn product(A, n/2) );  
    mult( spawn product(A+n/2, n-n/2) );  
    sync;  
    return p;  
}
```

Executing Cilk programs

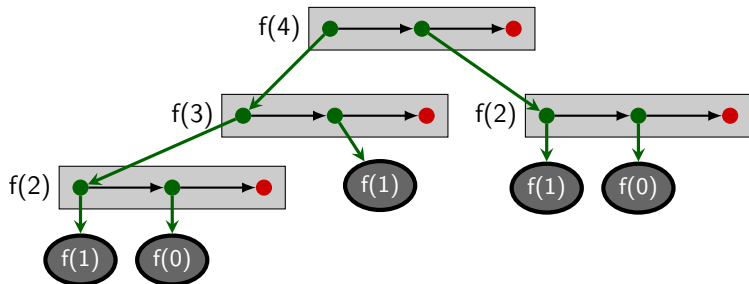
A dynamic graph (tasks, spawns, syncs)

```
cilk int fib(int n) {  
    if (n < 2) return (n);  
    x = spawn fib(n - 1);  
    y = spawn fib(n - 2);  
    sync;  
    return x + y;  
}
```

Executing Cilk programs

A dynamic graph (tasks, spawns, syncs)


```
cilk int fib(int n) {  
    if (n < 2) return (n);  
    x = spawn fib(n - 1);  
    y = spawn fib(n - 2);  
    sync;  
    return x + y;  
}
```



Executing Cilk programs

A dynamic graph (tasks, spawns, syncs)

```
cilk int fib(int n) {  
    if (n < 2) return (n);  
    x = spawn fib(n - 1);  
    y = spawn fib(n - 2);  
    sync;  
    return x + y;  
}
```



How do we efficiently execute Cilk programs?

The work-first principle: Minimize the scheduling overhead borne by the work of a computation. Specifically, move overheads out of the work and onto the critical path.


(i.e., make common case fast, push overheads to the rare case)

How do we efficiently execute Cilk programs?

The work-first principle: Minimize the scheduling overhead borne by the **work** of a computation. Specifically, move overheads out of the **work** and onto the **critical path**.


(i.e., make common case fast, push overheads to the rare case)

Performance model




- T_p : Execution time on P CPUs

Performance model




- T_p : Execution time on P CPUs
- T_1 : **work**
(Execution time for all nodes)

Performance model




- T_p : Execution time on P CPUs
- T_1 : **work**
(Execution time for all nodes)
- T_∞ : **span** / **critical path**
(Execution time for ∞ CPUs)


Scheduling tasks




P_0



P_1




P_2




P_3

Scheduling tasks



work sharing: when new tasks are created, scheduler tries to send them to inactive CPUs


Scheduling tasks



work stealing: Idle processors try to steal tasks

- ▶ child executes first
- ▶ tasks are stolen from the bottom
- ▶ intuitively similar to DFS (good space properties)
- ▶ where do we steal from?

Scheduling tasks



work stealing: Idle processors try to steal tasks

- ▶ child executes first
- ▶ tasks are stolen from the bottom
- ▶ intuitively similar to DFS (good space properties)
- ▶ where do we steal from? **randomly choose**

Cilk's work-stealing scheduler

- ▶ $T_p = T_1/P + \mathcal{O}(T_\infty)$
 - (T_1/P and T_∞ are lower bounds)
- ▶ requires $S_p \leq PS_1$ stack space

see: Scheduling Multithreaded Computations by Work Stealing, by Blumofe and Leiserson.

Parallel slackness

Cilk's basic assumption

- ▶ program parallelism T_1/T_∞
- ▶ machine parallelism P
- ▶ $P \ll T_1/T_\infty$
- ▶ parallel slack: $(T_1/T_\infty)/P$

$$\left. \begin{array}{l} T_p = T_1/P + \mathcal{O}(T_\infty) \\ T_1/P \gg T_\infty \end{array} \right\} \Rightarrow T_p \approx T_1/P \quad (\text{linear speedup})$$

- ▶ Intuitively:
 - + program does not depend on number of CPUs
 - + allows better load balancing

Parallel slackness \Rightarrow work-first principle

$$T_p \approx T_1/P$$
$$c_1 = T_1/T_s \Rightarrow T_p \approx c_1 T_s/P \quad T_s: \text{time of C elision}$$

principle: minimize c_1 even at the expense of c_∞
(yet, this has limits – authors' discussion on Cilk-4)

Compilation strategy

Two versions of each Cilk function (task)


- ▶ fast version
 - ▶ invariant: never stolen
 - ▶ as close to C elision as possible
- ▶ slow version
 - ▶ when a task is stolen, slow version is executed on thief processor
 - ▶ needs to restore state
- ▶ Original implementation: cilk2c translation

C elision

Only use the execution stack

```
f(n) {  
    // spawn f(n-1)  
    x = f(n-1);  
    // spawn f(n-2)  
    y = f(n-2);  
    // sync  
    return x + y;  
}
```

Cactus stack



Cilk fast version

use the execution stack, record state

```
f(n) {  
    frame = alloc_frame();  
    // spawn f(n-1)  
    frame->entry = 1;  
    frame->n = n;  
    push(frame);  
    x = f(n-1);  
    frame = pop();  
    if (!frame)  
        return STOLEN;  
    // spawn f(n-2)  
    ...  
    // sync  
    free_frame(frame)  
    return x + y;  
}
```

Cilk slow version

restore state, sync

```
f_slow(frame) {  
    switch (frame->entry) {  
        case 1: goto a;  
        ...  
    }  
    // spawn f(n-1)  
    ... (same as fast version) ...  
    if (0) { a: n = frame->n; }  
    // spawn f(n-2)  
    ... (same as fast version) ...  
    // sync  
    wait_for_children();  
    free_frame(frame)  
    return frame->x + frame->y;  
}
```

worker/thief synchronization

- ▶ basic scheduler data structure: double-ended queue (*deque*)
- ▶ worker: operates on **T**ail: push / pop
- ▶ thief: operates on **H**ead: steal
- ▶ traditional approach: lock for accessing deque
 - ▶ adds overhead to worker!

TH(E) protocol

<i>Worker/Victim</i>	<i>Thief</i>
1 push() {	1 steal() {
2 T++;	2 lock(L);
3 }	3 H++;
	4 if (H > T) {
4 pop() {	5 H--;
5 T--;	6 unlock(L);
6 if (H > T) {	7 return FAILURE;
7 T++;	8 }
8 lock(L);	9 unlock(L);
9 T--;	10 return SUCCESS;
10 if (H > T) {	11 }
11 T++;	
12 unlock(L);	
13 return FAILURE;	
14 }	
15 unlock(L);	
16 }	
17 return SUCCESS;	
18 }	

Note #1: improving performance

```
def rle_rec(xs):  
    if len(xs) <= cutoff:  
        return rle(xs)  
    mid = len(xs) // 2  
    rle1 = spawn rle_rec(xs[mid:])  
    rle2 = spawn rle_rec(xs[:mid])  
    sync  
    return rle_merge(rle1, rle2)
```

Note #2: data structures

- ▶ efficient partition and concatation
- ▶ lists: poor partition
- ▶ arrays: poor concatation
- ▶ Ropes: an Alternative to Strings
Boehm et al. (1995)
- ▶ Skip lists: a probabilistic alternative to balanced trees
Pugh (1990)
- ▶ implementation: <https://github.com/kkourt/xarray>

Note #3: data vs task parallelism

```
reduce(rle_merge, map(lambda x: [(x,1)], input))
```

(rle_merge is an associative operation)

Thank you!

Qs? Pizza?